

PREDICTING H1N1 FLU VACCINATION STATUS USING MACHINE LEARNING

1. Business Understanding

1.1 Overview

In this project, the aim was to use data from the National Flu Survey (NHFS 2009) to predict whether respondents received the H1N1 vaccine. Understanding past vaccination trends is crucial for interpreting patterns in more recent pandemics, such as COVID-19. Key factors influencing vaccination status include Doctor recommendations for the H1N1 vaccine, health insurance, opinions on the vaccine's effectiveness and perceptions of the risk posed by H1N1. I employed six machine learning models for prediction:

1. Decision Tree Classifier
2. Logistic Regression
3. Random Forest
4. K-Nearest Neighbors Classifier
5. Gradient Boosting Classifier
6. XGBoost Classifier

Among these, the Gradient Boosting Classifier achieved the highest accuracy and precision.

1.2 Business Problem.

Vaccination stands as one of the most effective public health interventions ever implemented, leading to the elimination and control of diseases that were once widespread globally. Despite substantial medical evidence and the strong consensus among healthcare professionals supporting vaccination, skepticism has increased in many countries in recent years. This troubling trend has resulted in decreased immunization coverage, with several outbreaks of infectious diseases linked to undervaccinated communities. The growing issue of vaccine hesitancy has become so pervasive that it is now the subject of numerous studies aiming to understand the sources and correlations of attitudes toward vaccination.

This study aims to predict the likelihood of individuals receiving the H1N1 flu vaccine. We believe the predictive models and analyses from this study will provide public health professionals and policymakers with a clear understanding of the factors associated with low vaccination rates. This, in turn, will enable them to systematically address the barriers preventing people from getting vaccinated.

The methodologies employed in these models can serve as a reference for future work and can be compared with other models for performance evaluation. To accurately classify those who received the H1N1 flu shot from those who did not, we require models with high accuracy and high precision, which corresponds to a low false positive rate (those mistakenly identified as vaccinated when they were not). This will be further evaluated using the ROC curve, accuracy score, precision score, and confusion matrix.

Target Audience: Public health officials of the American Public Health Association (APHA)

OBJECTIVES:

1. Predicting who is vaccinated or not accurately.(Deliverable: Model)
2. Analyse the factors that influence people to get H1N1 vaccine or not. (Deliverable: Analysis)

Context:

- False negative: Saying people did not get the vaccine when they actually did.
- Outcome: Not a big problem
- False positive: Saying people got the vaccine when they actually did not.
- Outcome: Big problem

Evaluation: We will focus on accuracy, f1, and precision scores for our model iterations in order to minimize False Positives, because in our business context false positives are a much more costly mistake than false negatives.

- **Accuracy**
- **Precision**
- Recall
- **F1-Score**

2. Data Understanding

2.1 Importing the necessary libraries and exploring the data

```
In [ ]: # Common libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Libraries for model training
```

```

from sklearn.preprocessing import StandardScaler, MinMaxScaler, MaxAbsScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
import category_encoders as ce
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score

# Libraries for algorithm

from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
import xgboost      # extreme gradient boosting

# Libraries for testing
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score, roc_auc_score,
    ConfusionMatrixDisplay, confusion_matrix
)

# Removing warnings
import warnings
warnings.filterwarnings('ignore')

# Storing plots

%matplotlib inline

# To visualize the 100 many the columns in data
pd.options.display.max_columns=100

```

2.2 Load Dataset

This data comes from a NHFS National Flu Survey from 2009, which inquires about whether or not people received the seasonal flu and/or the H1N1 flu vaccination, as well as their demographic, behavioral, and health factors.

```

In [ ]: # Reading in the data and previewing the dataset
Data1 = pd.read_csv('DATA/H1N1_Flu_Vaccines.csv')
Data1.head(5)

```

Out[]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavior
0	0	1.0	0.0	0.0	
1	1	3.0	2.0	0.0	
2	2	1.0	1.0	0.0	
3	3	1.0	1.0	0.0	
4	4	2.0	1.0	0.0	

In []: *#dataset tail*
 Data1.tail(3)

Out[]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavior
26704	26704	2.0	2.0	0.0	
26705	26705	1.0	1.0	0.0	
26706	26706	0.0	0.0	0.0	

2.3 Checking the Dataset

In []: *#Determining the no. of records in our dataset*
 Data1.shape

Out[]: (26707, 38)

In []: *# Exploring the percentage breakdown of the two classes in one possible t*
 Data1['seasonal_vaccine'].value_counts(normalize=True)

Out[]: seasonal_vaccine
 0 0.534392
 1 0.465608
 Name: proportion, dtype: float64

In []: *# Exploring the percentage breakdown of the two classes in one possible t*
 Data1['h1n1_vaccine'].value_counts(normalize=True) *# class imbalance pr*

Out[]: h1n1_vaccine
 0 0.787546
 1 0.212454
 Name: proportion, dtype: float64

```
In [ ]: # checking dataset information
        Data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 38 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   respondent_id                        26707 non-null   int64
1   hlnl_concern                        26615 non-null   float64
2   hlnl_knowledge                      26591 non-null   float64
3   behavioral_antiviral_meds          26636 non-null   float64
4   behavioral_avoidance                26499 non-null   float64
5   behavioral_face_mask                26688 non-null   float64
6   behavioral_wash_hands               26665 non-null   float64
7   behavioral_large_gatherings         26620 non-null   float64
8   behavioral_outside_home              26625 non-null   float64
9   behavioral_touch_face                26579 non-null   float64
10  doctor_recc_hlnl                    24547 non-null   float64
11  doctor_recc_seasonal                 24547 non-null   float64
12  chronic_med_condition                25736 non-null   float64
13  child_under_6_months                 25887 non-null   float64
14  health_worker                       25903 non-null   float64
15  health_insurance                     14433 non-null   float64
16  opinion_hlnl_vacc_effective            26316 non-null   float64
17  opinion_hlnl_risk                      26319 non-null   float64
18  opinion_hlnl_sick_from_vacc            26312 non-null   float64
19  opinion_seas_vacc_effective            26245 non-null   float64
20  opinion_seas_risk                      26193 non-null   float64
21  opinion_seas_sick_from_vacc            26170 non-null   float64
22  age_group                            26707 non-null   object
23  education                            25300 non-null   object
24  race                                 26707 non-null   object
25  sex                                  26707 non-null   object
26  income_poverty                       22284 non-null   object
27  marital_status                       25299 non-null   object
28  rent_or_own                           24665 non-null   object
29  employment_status                    25244 non-null   object
30  hhs_geo_region                       26707 non-null   object
31  census_msa                           26707 non-null   object
32  household_adults                     26458 non-null   float64
33  household_children                   26458 non-null   float64
34  employment_industry                  13377 non-null   object
35  employment_occupation                13237 non-null   object
36  hlnl_vaccine                         26707 non-null   int64
37  seasonal_vaccine                     26707 non-null   int64
dtypes: float64(23), int64(3), object(12)
memory usage: 7.7+ MB
```

```
In [ ]: Data1.dtypes
```

```
Out[ ]: respondent_id      int64
        h1n1_concern      float64
        h1n1_knowledge     float64
        behavioral_antiviral_meds float64
        behavioral_avoidance float64
        behavioral_face_mask float64
        behavioral_wash_hands float64
        behavioral_large_gatherings float64
        behavioral_outside_home float64
        behavioral_touch_face float64
        doctor_recc_h1n1    float64
        doctor_recc_seasonal float64
        chronic_med_condition float64
        child_under_6_months float64
        health_worker       float64
        health_insurance    float64
        opinion_h1n1_vacc_effective float64
        opinion_h1n1_risk    float64
        opinion_h1n1_sick_from_vacc float64
        opinion_seas_vacc_effective float64
        opinion_seas_risk    float64
        opinion_seas_sick_from_vacc float64
        age_group           object
        education           object
        race                object
        sex                 object
        income_poverty      object
        marital_status      object
        rent_or_own         object
        employment_status   object
        hhs_geo_region      object
        census_msa          object
        household_adults    float64
        household_children  float64
        employment_industry object
        employment_occupation object
        h1n1_vaccine        int64
        seasonal_vaccine    int64
        dtype: object
```

```
In [ ]: # Getting number of null values.
        Data1.isna().sum()
```

```
Out[ ]: respondent_id      0
        h1n1_concern      92
        h1n1_knowledge    116
        behavioral_antiviral_meds  71
        behavioral_avoidance  208
        behavioral_face_mask  19
        behavioral_wash_hands  42
        behavioral_large_gatherings  87
        behavioral_outside_home  82
        behavioral_touch_face  128
        doctor_recc_h1n1    2160
        doctor_recc_seasonal  2160
        chronic_med_condition  971
        child_under_6_months  820
        health_worker      804
        health_insurance    12274
        opinion_h1n1_vacc_effective  391
        opinion_h1n1_risk    388
        opinion_h1n1_sick_from_vacc  395
        opinion_seas_vacc_effective  462
        opinion_seas_risk    514
        opinion_seas_sick_from_vacc  537
        age_group          0
        education          1407
        race               0
        sex               0
        income_poverty     4423
        marital_status     1408
        rent_or_own        2042
        employment_status  1463
        hhs_geo_region     0
        census_msa         0
        household_adults   249
        household_children  249
        employment_industry 13330
        employment_occupation 13470
        h1n1_vaccine       0
        seasonal_vaccine   0
        dtype: int64
```

```
In [ ]: Data1.duplicated().sum()
```

```
Out[ ]: 0
```

```
In [ ]: # Explore numerical columns
        Data1.describe()
```

Out[]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	beh
count	26707.000000	26615.000000	26591.000000		26636.000000
mean	13353.000000	1.618486	1.262532		0.048844
std	7709.791156	0.910311	0.618149		0.215545
min	0.000000	0.000000	0.000000		0.000000
25%	6676.500000	1.000000	1.000000		0.000000
50%	13353.000000	2.000000	1.000000		0.000000
75%	20029.500000	2.000000	2.000000		0.000000
max	26706.000000	3.000000	2.000000		1.000000

In []: *# Explore object columns*
 Data1[[c for c in Data1.columns if Data1[c].dtype == 'object']].describe()

Out[]:

	age_group	education	race	sex	income_poverty	marital_status	rent_o
count	26707	25300	26707	26707	22284		25299
unique	5	4	4	2	3		2
top	65+ Years	College Graduate	White	Female	<= \$75,000, Above Poverty		Married
freq	6843	10097	21222	15858	12777		13555

Observations

Upon initial exploration of the data, we've made several key observations:

1. There are 26,000 respondents to this survey.
2. There are 36 features.
3. There are lots of missing value so we need to impute them.
4. There no duplicates in dataset.

Further preprocessing is required to understand the relationships between different features.

I decided to choose the H1N1 vaccination rate as our target variable, because so many of the features are related to H1N1 vaccination.

3. EXPLORATORY DATA ANALYSIS(EDA)

How many people got the H1N1 vaccine - from the given data set


In []: *## for the H1N1 vaccine*
 fig1, ax1 = plt.subplots()


```

labels = ['No', 'Yes']
explode = (0, 0.1)
colors = ['#ff9999', '#66b3ff']
ax1.pie(Data1['h1n1_vaccine'].value_counts(), explode=explode, labels=labels,
        shadow=True, startangle=90)
plt.title("Less than 25% people received the H1N1 vaccine")
ax1.axis('equal')
plt.tight_layout()
plt.savefig("H1n1pie")
plt.show()

```

Less than 25% people received the H1N1 vaccine



Observation - There is a class imbalance in the adoption of the H1N1 vaccine. (less than 25% of the people chose to receive the H1N1 vaccine). This might help in future during the analysis. This class imbalance problem is what we want to deal with in this project.

The following visualizations represent the top four most influential features in determining vaccination status for H1N1.

```

In [ ]: # Making a copy of main dataframe to use for visualizations
Data2 = Data1.copy()

```

```

In [ ]: # Creating dictionary for mapping in order to create better names for x axis
ins_dict = {1: 'Health Insurance', 0: 'No Health Insurance'}
# Creating the column that will be used to create clear x axis tick marks
Data2['health_ins_words'] = Data2['health_insurance'].replace(ins_dict)

```

```

In [ ]: def plot_bar(df, x_col, y_col, ax, xlabel, ylabel, title):
    """
    This function plots a bar chart on a given axis.
    """
    sns.barplot(x=df[x_col].dropna(), y=df[y_col]*100, ax=ax)

```

```

ax.set_xlabel(xlabel)
ax.set_ylabel(ylabel, fontsize=13)
ax.set_title(title, fontsize=13)

def plot_vaccination_graphs(Data2):
    """
    This function takes a dataframe and plots four bar charts in a 2x2 layout
    showing the relationship between H1N1 vaccination and various factors
    """
    # Making a copy of the main dataframe to use for visualizations
    Data2 = Data1.copy()

    # Creating dictionary for mapping in order to create better names for
    ins_dict = {1: 'Health Insurance', 0: 'No Health Insurance'}
    # Creating the column that will be used to create clear x axis tick marks
    Data2['health_ins_words'] = Data2['health_insurance'].replace(ins_dict)

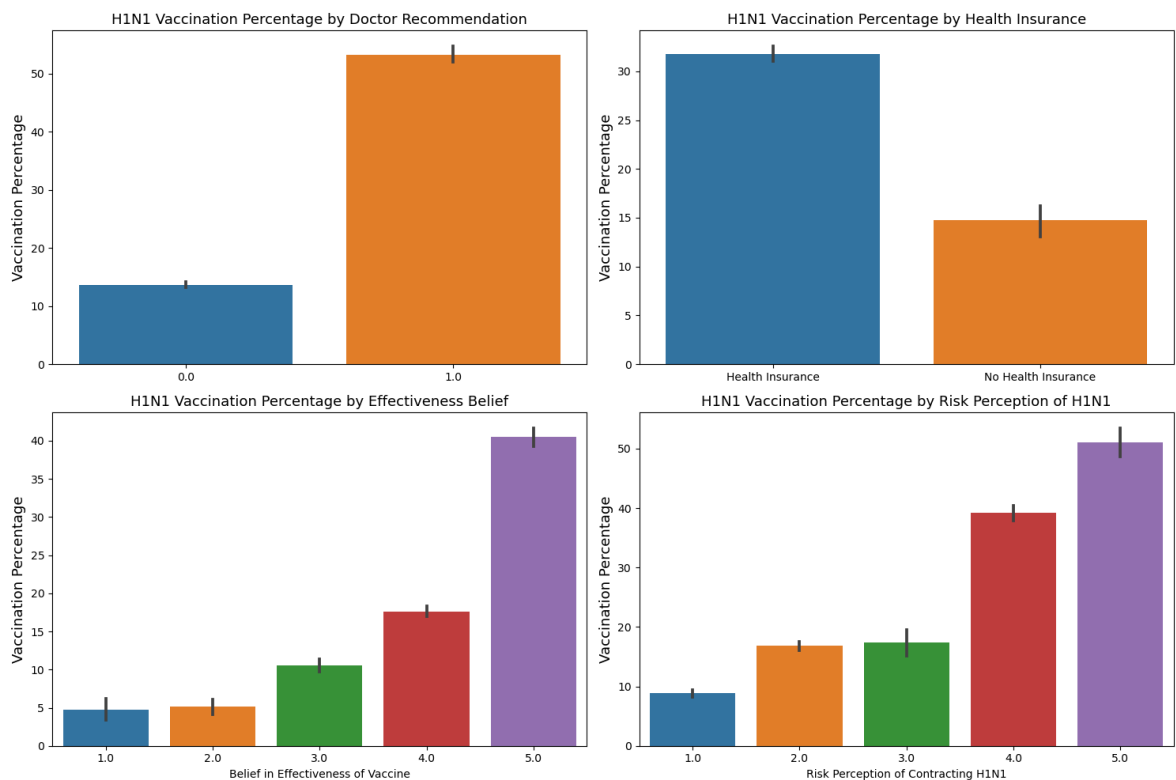
    # Setting up the 2x2 subplot layout
    fig, axs = plt.subplots(2, 2, figsize=(15, 10))

    # Plotting the individual bar charts
    plot_bar(Data2, 'doctor_recc_h1n1', 'h1n1_vaccine', axs[0, 0], '', 'V')
    plot_bar(Data2, 'health_ins_words', 'h1n1_vaccine', axs[0, 1], '', 'V')
    plot_bar(Data2, 'opinion_h1n1_vacc_effective', 'h1n1_vaccine', axs[1, 0], 'V')
    plot_bar(Data2, 'opinion_h1n1_risk', 'h1n1_vaccine', axs[1, 1], 'Risk')

    # Adjust layout
    plt.tight_layout()
    plt.show()

# Call the function to plot the graphs
plot_vaccination_graphs(Data2)

```



Observations:-

1. The plot "H1N1 Vaccination Percentage by Doctor Recommendation" shows a higher vaccination percentage among individuals who received a doctor's recommendation for the H1N1 vaccine compared to those who did not
2. The plot "H1N1 Vaccination Percentage by Health Insurance" might reveals a higher vaccination percentage among individuals with health insurance compared to those without it. This observation indicates that having health insurance could positively impact an individual's likelihood of getting vaccinated, possibly due to better access to healthcare services.
3. The plot "H1N1 Vaccination Percentage by Effectiveness Belief" probably indicates that individuals who believe in the effectiveness of the H1N1 vaccine have a higher vaccination percentage.
4. The plot "H1N1 Vaccination Percentage by Risk Perception of H1N1" is expected to show that individuals who perceive a higher risk of contracting H1N1 have a higher vaccination percentage.

```
In [ ]: # Function to find the outliers

def findoutliers(column):
    outliers=[]
    Q1=column.quantile(.25)
    Q3=column.quantile(.75)
    IQR=Q3-Q1
    lower_limit=Q1-(1.5*IQR)
    upper_limit=Q3+(1.5*IQR)
    for out1 in column:
        if out1>upper_limit or out1 <lower_limit:
            outliers.append(out1)


    return np.array(outliers)
```

```
In [ ]: print(len(findoutliers(Data2.household_adults)))
        print(len(findoutliers(Data2.household_children)))
```

```
1125
1747
```


```
In [ ]: # Visualising the outliers
        sns.boxplot(Data2.household_adults)
```

```
Out[ ]: <Axes: >
```



```
In [ ]: sns.boxplot(Data2.household_children)
```

```
Out[ ]: <Axes: >
```



There are outliers in the dataset but we are not removing them as some algorithms are not sensitive to outliers

4. Data Preparation (Data Cleaning)

There were a few changes I made to the data set.

1. First, I dropped the "respondent_id" and "seasonal_vaccine" columns because they were not relevant for our purposes.
2. I also added columns due to categorical columns that we transformed with OneHotEncoder.
3. I also filled null values with Iterative Imputer, which was a better alternative to simple imputer for our dataset.
4. I replaced category names with frequency counts with CountEncoder for the columns which had more than 10 unique categories.
5. I used pipelines to make preprocessing and modelling more efficient, and also to prevent data leakage.
6. I also decided to split training and testing data twice so that we could have a holdout set to test our final model's generalizability at the end.

```
In [ ]: # Define our X and y
X = Data1.drop(columns = ['respondent_id', 'h1n1_vaccine', 'seasonal_vaccine'])
y = Data1['h1n1_vaccine']
```

I chose 80%, 20% for train and validation. Also set the random seed to 42 and stratify=y.

Using stratify=y ensures that the holdout set has a similar distribution of classes as the original dataset, which is particularly important when dealing with imbalanced datasets to ensure that the model's evaluation is fair and representative.

4.1 Train-Test Split

I decided to split training and testing data twice so that we could have a holdout set to test our final model's generalizability at the end.


```
In [ ]: # Train - Holdout Set Split
X_train, X_hold, y_train, y_hold = train_test_split(X, y, test_size=0.2,
```

```
In [ ]: X_train.head()
```

Out []:

	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance
--	--------------	----------------	---------------------------	----------------------

20417	1.0	2.0	0.0	1.0
13969	2.0	2.0	0.0	1.0
24930	2.0	2.0	0.0	1.0
15420	2.0	1.0	0.0	0.0
10998	2.0	1.0	0.0	0.0



We used the stratify argument for y (our target) in both splits to help deal with the class imbalance problem.

```
In [ ]: # Regular Train Test Split
X_tr, X_te, y_tr, y_te = train_test_split(X_train, y_train, test_size=0.2)
```

```
In [ ]: # Set up lists for each columns datatypes
num_cols = []
ohe_cols = []
freq_cols = []

for c in X.columns:
    if X[c].dtype in ['float64', 'int64']:
        num_cols.append(c)
    elif X[c].nunique() < 10:
        ohe_cols.append(c)
    else:
        freq_cols.append(c)
```

```
In [ ]: # We wanted to see each column category
print(f'Numerical Columns:', num_cols)
print('\n')
print(f'Object Columns (with less than 10 unique values):', ohe_cols)
print('\n')
print(f'Object Columns (with more than 10 unique values):', freq_cols)
```

Numerical Columns: ['h1n1_concern', 'h1n1_knowledge', 'behavioral_antiviral_meds', 'behavioral_avoidance', 'behavioral_face_mask', 'behavioral_wash_hands', 'behavioral_large_gatherings', 'behavioral_outside_home', 'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal', 'chronic_med_condition', 'child_under_6_months', 'health_worker', 'health_insurance', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk', 'opinion_h1n1_sick_from_vacc', 'opinion_seas_vacc_effective', 'opinion_seas_risk', 'opinion_seas_sick_from_vacc', 'household_adults', 'household_children']

Object Columns (with less than 10 unique values): ['age_group', 'education', 'race', 'sex', 'income_poverty', 'marital_status', 'rent_or_own', 'employment_status', 'census_msa']

Object Columns (with more than 10 unique values): ['hhs_geo_region', 'employment_industry', 'employment_occupation']

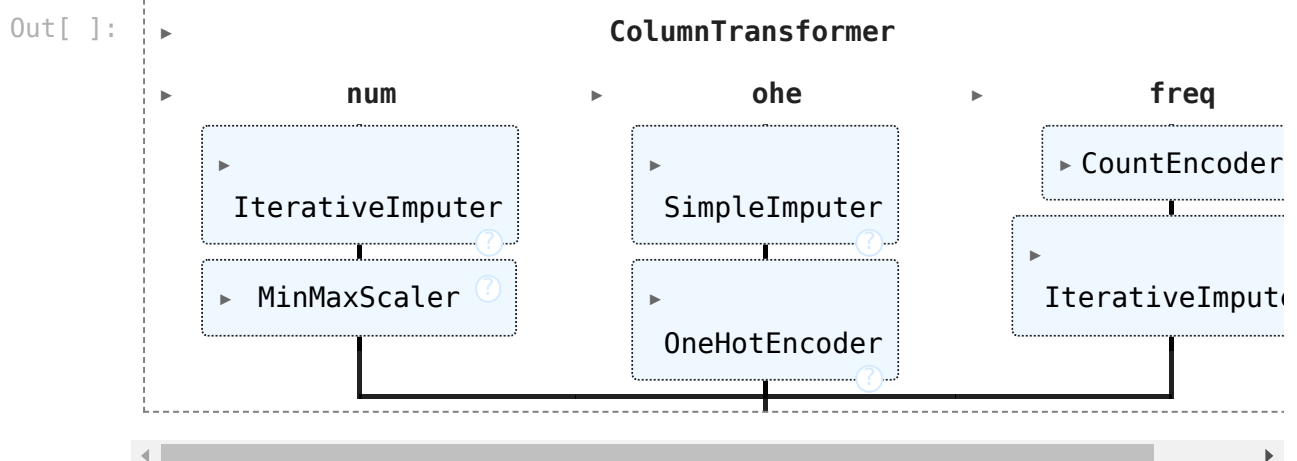
```
In [ ]: # Preprocessing with Pipelines
num_transformer = Pipeline(steps=[
    ('num_imputer', IterativeImputer(max_iter=100, random_state=42)), #
    ('minmaxscaler', MinMaxScaler()) #
])

ohe_transformer = Pipeline(steps=[
    ('ohe_imputer', SimpleImputer(strategy='constant', fill_value='Unknown')),
    ('ohe_encoder', OneHotEncoder(handle_unknown='ignore'))
])

freq_transformer = Pipeline(steps=[
    ('freq_encoder', ce.count.CountEncoder(normalize=True, min_group_size=10)),
    ('freq_imputer', IterativeImputer(max_iter=100, random_state=42))
])
```

```
In [ ]: # Preprocessor defined using ColumnTransformer by packaging the all components
preprocessor = ColumnTransformer(
    transformers=[
        ('num', num_transformer, num_cols),
        ('ohe', ohe_transformer, ohe_cols),
        ('freq', freq_transformer, freq_cols)
    ]
)
```

```
In [ ]: # Fitting preprocessor to see the components as a whole to the training set
preprocessor.fit(X_tr)
```



```
In [ ]: # Let's see what this looks like after the preprocessor transformation
X_tr_transformed = preprocessor.transform(X_tr)
X_tr_transformed.shape
```

```
Out[ ]: (17092, 59)
```

The number of features increase from 36 to 59 while the no of rows reduce from 26707 to 17092.

```
In [ ]: # Visualize it with Pandas dataframe
pd.DataFrame(X_tr_transformed).head()
```

```
Out[ ]:
```

	0	1	2	3	4	5	6	7	8	9
0	1.000000	0.5	0.008339	0.000000	1.000000	0.884731	1.0	1.000000	1.0	1.000000
1	0.666667	1.0	0.008339	0.966636	0.038403	0.884731	0.0	0.020289	1.0	1.000000
2	0.000000	0.5	0.008339	0.966636	0.038403	0.884731	0.0	0.020289	1.0	1.000000
3	0.666667	1.0	0.008339	0.966636	0.038403	0.884731	0.0	0.020289	1.0	0.137273
4	0.666667	1.0	0.008339	0.966636	0.038403	0.884731	1.0	1.000000	1.0	0.137273

Modeling

I wanted to use a variety of different models so as to find the most accurate model.

Because there are many different hyperparameters for each model and I did not know the optimal combinations, we used GridSearchCV to find the best combinations for each model. I specified class weight to be balanced in order to address the class imbalance issue for our models, whenever possible. I also analyzed the accuracy score, precision score, f1 score, and roc-auc curve for each model.

I also compared the different roc-auc curves of each model, to choose the final model. Additionally, I looked closely at the confusion matrix to see whether or not we were minimizing false positives. Gradient Boosting Classifier gave us the best accuracy and precision scores, so we chose it to be our final model.

We will use this function to evaluate the performance of our models:

```
In [ ]: def evaluate(estimator, X_tr, X_te, y_tr, y_te, roc_auc='skip'):
        """
        Evaluation function to show a few scores for both the train and test
        Also shows a confusion matrix for the test set.

        Parameters:
        estimator : a fit sklearn-style model or pipeline
        X_tr : array or pandas DataFrame
                Training input variables
        X_te : array or pandas DataFrame
                Testing input variables
        y_tr : array or pandas Series
                Training output variable
```



```

y_te : array or pandas Series
    Testing output variable
roc_auc : str
    'skip': default, skips calculating roc_auc
    'dec': use decision_function to calculate roc_auc
    'proba': use predict_proba to calculate roc_auc
"""

# Grab predictions
tr_preds = estimator.predict(X_tr)
te_preds = estimator.predict(X_te)

# Output needed for roc_auc_score
if roc_auc == 'skip': # skips calculating the roc_auc_score
    train_out = False
    test_out = False
elif roc_auc == 'dec':
    train_out = estimator.decision_function(X_tr)
    test_out = estimator.decision_function(X_te)
elif roc_auc == 'proba':
    train_out = estimator.predict_proba(X_tr)[: , 1] # proba for the 1
    test_out = estimator.predict_proba(X_te)[: , 1]
else:
    raise ValueError("The value for roc_auc should be 'skip', 'dec' or 'proba'")

# Print training scores
print("Training Scores:")
print(f"Train Accuracy: {accuracy_score(y_tr, tr_preds)}")
print(f"Train Precision: {precision_score(y_tr, tr_preds)}")
print(f"Train Recall: {recall_score(y_tr, tr_preds)}")
print(f"Train F1-Score: {f1_score(y_tr, tr_preds)}")
if isinstance(train_out, np.ndarray): # checking for roc_auc
    print(f"Train ROC-AUC: {roc_auc_score(y_tr, train_out)}")

print("\n" * 10)

# Print testing scores
print("Testing Scores:")
print(f"Test Accuracy: {accuracy_score(y_te, te_preds)}")
print(f"Test Precision: {precision_score(y_te, te_preds)}")
print(f"Test Recall: {recall_score(y_te, te_preds)}")
print(f"Test F1-Score: {f1_score(y_te, te_preds)}")
if isinstance(test_out, np.ndarray): # checking for roc_auc
    print(f"Test ROC-AUC: {roc_auc_score(y_te, test_out)}")

# Plot confusion matrix for test set
ConfusionMatrixDisplay.from_estimator(estimator, X_te, y_te, cmap="pl
plt.show()

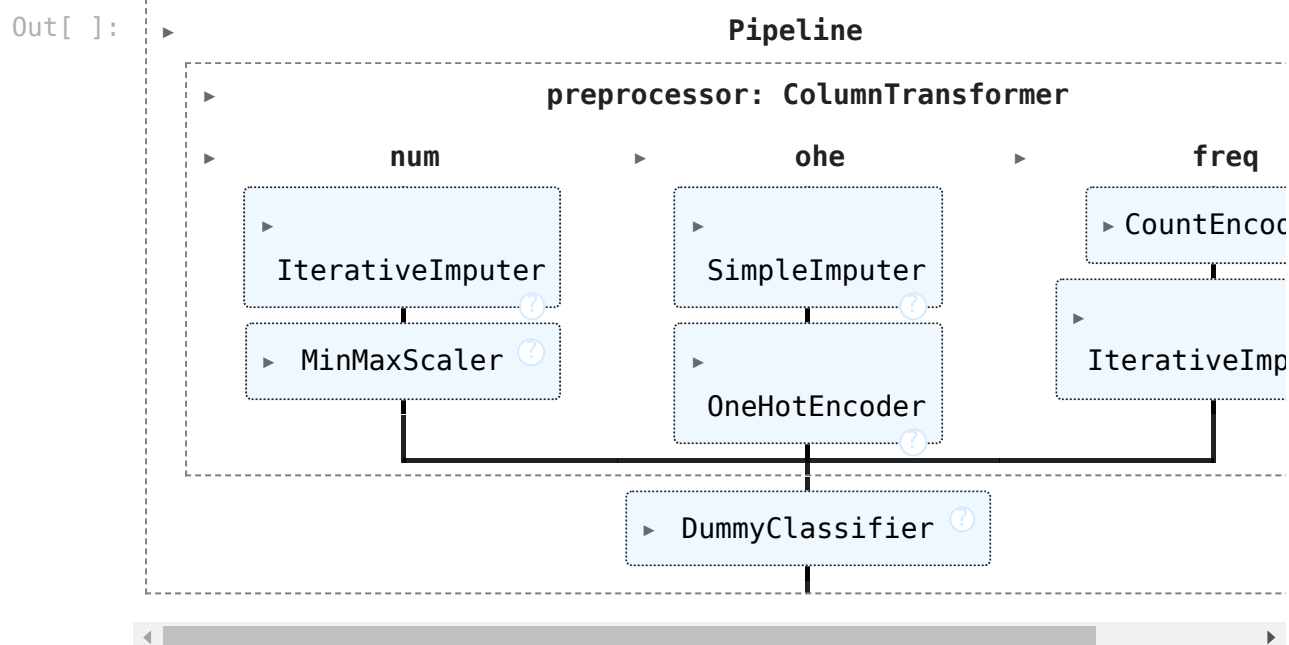
```

Baseline Understanding

- To be able to truly understand and then improve our model's performance, we first need to establish a baseline for the data that we have
- Let's use DummyClassifier to make prediction based on the most frequent class in the target variable, which is 0 in our case.

```
In [ ]: # Setting the up the dummy model to go through the pipeline
dummy_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', DummyClassifier(strategy="most_frequent"))
])
```

```
In [ ]: # Fitting the dummy model
dummy_model.fit(X_tr, y_tr)
```



```
In [ ]: # Evaluate dummy model using the evaluate function created above
evaluate(dummy_model, X_tr, X_te, y_tr, y_te, roc_auc='skip')
```

Training Scores:

Train Accuracy: 0.7875614322490054

Train Precision: 0.0

Train Recall: 0.0

Train F1-Score: 0.0


Testing Scores:

Test Accuracy: 0.7875029253451907

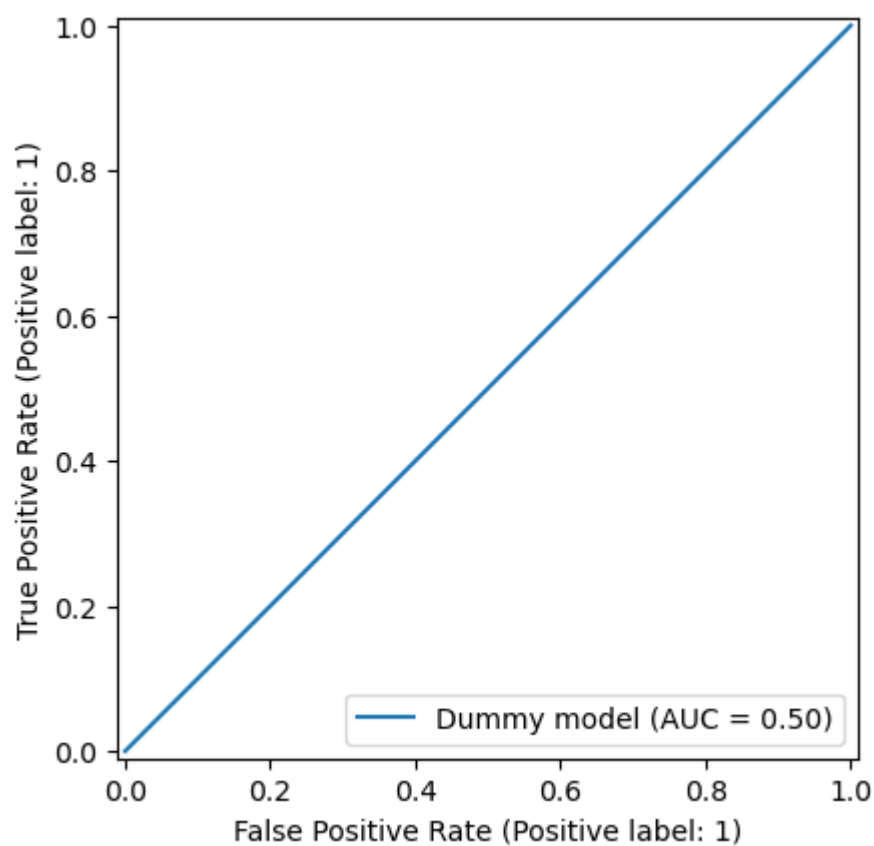
Test Precision: 0.0

Test Recall: 0.0

Test F1-Score: 0.0



```
In [ ]: # Plotting the ROC-AUC curve for the dummy model using from_estimator met
RocCurveDisplay.from_estimator(dummy_model, X_te, y_te, name='Dummy model')
plt.show()
```



So, the mean of the accuracy score is a little over 78% if we always guess the majority class, which is 0 in this case.

Modeling Iterations

Now we will start to iterate over multiple models!

MODEL 1: Decision Tree Classifier

```
In [ ]: # Setting up the DecisionTreeClassifier to go through the pipeline
dtc = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', DecisionTreeClassifier())
])
```

```
In [ ]: cross_validate(dtc, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([7.42701888, 5.89880443, 5.00161982, 8.07086277, 3.17
024493]),
'score_time': array([0.18780661, 0.08220649, 0.12744069, 0.08681822, 0.
0756011 ]),
'test_score': array([0.77127815, 0.7739105 , 0.77589233, 0.77940316, 0.
78730252]),
'train_score': array([1., 1., 1., 1., 1.])}
```

We see that, we have overfitting problem with DecisionTreeClassifier()!

```
In [ ]: # Let's do GridSearchCV
param_grid = {
    "classifier__max_depth": [1, 2, 5],
    "classifier__min_samples_split": [2, 10],
    "classifier__class_weight": ['balanced', None] # we have class-imba
}

# Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(dtc,
                    param_grid,
                    scoring=['f1', 'precision'],
                    refit = 'f1')

# Fit GridSearchCV
output_dtc = grid.fit(X_tr, y_tr)

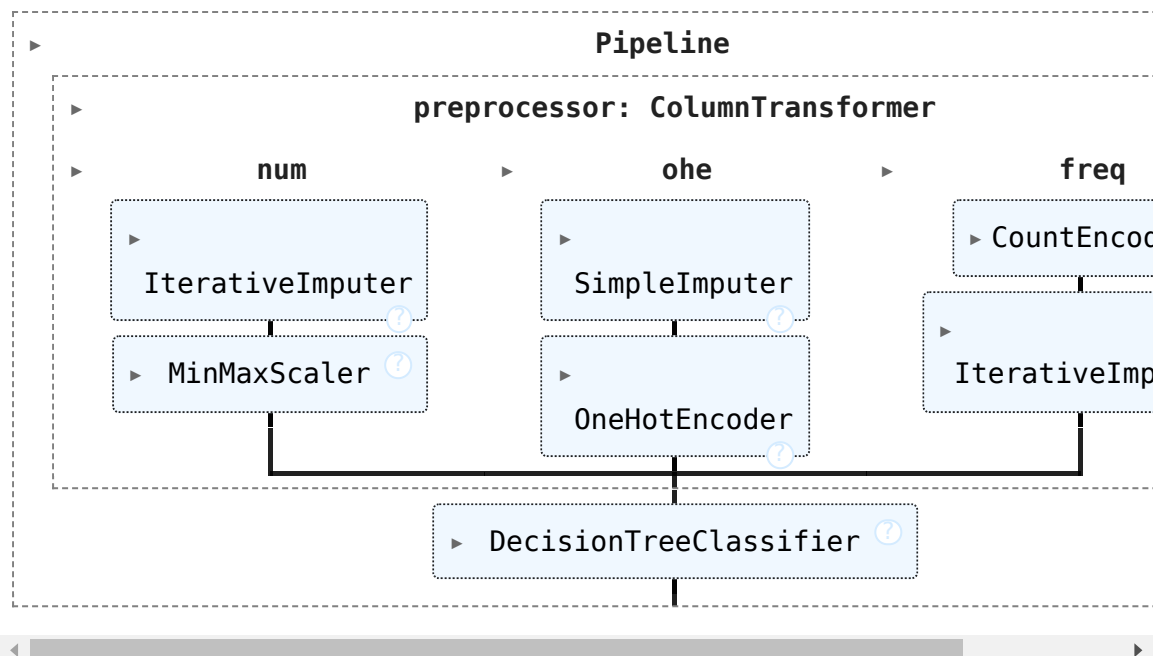
print(output_dtc.best_params_)

best_model_dtc=output_dtc.best_estimator_
```

```
{'classifier__class_weight': 'balanced', 'classifier__max_depth': 5, 'clas
sifier__min_samples_split': 10}
```

```
In [ ]: #Fit the model using the best parameters
best_model_dtc.fit(X_tr, y_tr)
```

Out[]:



```
In [ ]: # Evaluating the decision tree model for various metrics
        evaluate(best_model_dtc, X_tr, X_te, y_tr, y_te, roc_auc='skip')
```

Training Scores:

Train Accuracy: 0.794640767610578

Train Precision: 0.5115303983228512

Train Recall: 0.7391903057009088

Train F1-Score: 0.6046406848389276


Testing Scores:

Test Accuracy: 0.8045869412590686


Test Precision: 0.5279265493496557

Test Recall: 0.7599118942731278

Test F1-Score: 0.6230248306997742



```
In [ ]: # Plotting the ROC-AUC curve for the dummy model using from_estimator met
RocCurveDisplay.from_estimator(best_model_dtc, X_te, y_te, name='Decision
plt.show())
```



This decision tree model is not overfitting, but we have a low precision score, as well as a low f1 score. However, the AUC for this model is 0.85, which is fairly high, meaning that it does an adequate job of maximizing true positives and minimizing the false positives. This model is not overfitting.

MODEL 2: Logistic Regression

```
In [ ]: # Define your pipeline with preprocessing and classifier
logreg = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(max_iter=500, random_state=42))
])
```

```
In [ ]: cross_validate(logreg, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([4.93650317, 3.53602052, 2.53369737, 2.54796076, 2.62
224889]),
'score_time': array([0.07737327, 0.0631144 , 0.07117891, 0.0715673 , 0.
0787549 ]),
'test_score': array([0.83533197, 0.83503949, 0.82562902, 0.82709187, 0.
83762434]),
'train_score': array([0.83346742, 0.83346742, 0.83516162, 0.83618546,
0.83282141])}
```

```
In [ ]: # Define your grid for GridSearchCV
param_grid = {
```

```

"classifier__solver": ['lbfgs', 'liblinear', 'newton-cg', 'saga'],
"classifier__C" : [1, 0.1, 0.01], # regularization parameter
"classifier__class_weight": ['balanced', None]
}
# Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(logreg,
                    param_grid,
                    scoring=['f1', 'precision'],
                    refit = 'f1')

# Fit GridSearchCV
output_logreg = grid.fit(X_tr, y_tr)

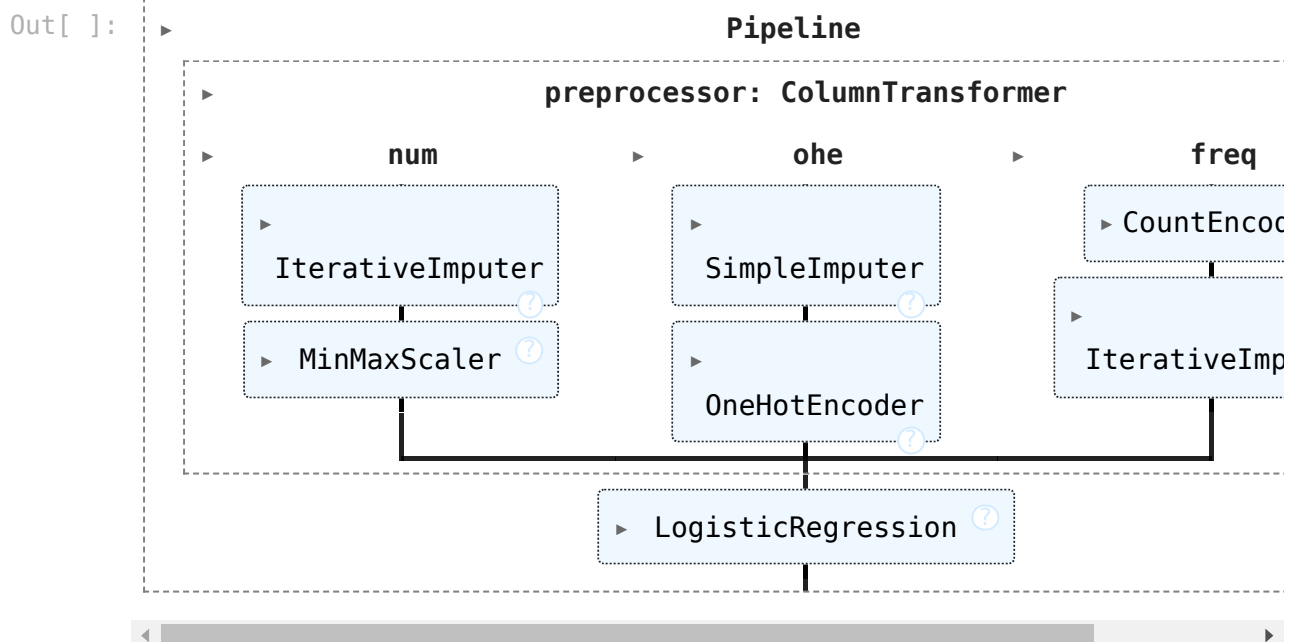
# print the best parameters for the model
print(output_logreg.best_params_)

# Retrieve the best estimator from GridSearchCV
best_model_logreg=output_logreg.best_estimator_

```

```
{'classifier__C': 1, 'classifier__class_weight': 'balanced', 'classifier__solver': 'saga'}
```

```
In [ ]: #Fit the best model to the training data
best_model_logreg.fit(X_tr, y_tr)
```



```
In [ ]: # Evaluating the logistic regression for various metrics
evaluate(best_model_logreg, X_tr, X_te, y_tr, y_te)
```

Training Scores:

```

Train Accuracy: 0.7753334893517435
Train Precision: 0.48130924700411376
Train Recall: 0.7411181492701735
Train F1-Score: 0.5836044242029929
*****


```

Testing Scores:


```

Test Accuracy: 0.7825883454247601
Test Precision: 0.4924187725631769
Test Recall: 0.751101321585903
Test F1-Score: 0.5948539031836023

```



```
In [ ]: RocCurveDisplay.from_estimator(best_model_logreg, X_te, y_te, name='logis  
plt.show()
```



This logistic regression model has low precision and fl scores, and has an AUC almost equal to the decision tree AUC above. This model is not overfitting.

MODEL 3: Random Forest

```
In [ ]: # Setting up the RandomForestClassifier to go through the pipeline
rfc = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', RandomForestClassifier(random_state=42))
])
```

```
In [ ]: cross_validate(rfc, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([5.17897797, 3.82426238, 3.97011614, 3.95122027, 4.16
435409]),
'score_time': array([0.12792277, 0.12009501, 0.12692761, 0.12853265, 0.
12861967]),
'test_score': array([0.85171103, 0.84907868, 0.84669397, 0.84698654, 0.
85225278]),
'train_score': array([1.          , 1.          , 1.          , 1.          ,
0.99985374])}
```

```
In [ ]: # Let's do GridSearchCV
param_grid = {
    "classifier__n_estimators": [100, 200],
    "classifier__max_depth" : [2, 5],
    "classifier__min_samples_leaf": [1, 2],
    "classifier__class_weight" : ['balanced', 'balanced_subsample']
    # class weight should be balanced or balanced_subsample - we have imbalanced
}

#Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(rfc,
                    param_grid,
                    scoring=['f1', 'precision'],
                    refit = 'f1')

# Fit GridSearchCV
output_rfc = grid.fit(X_tr, y_tr)

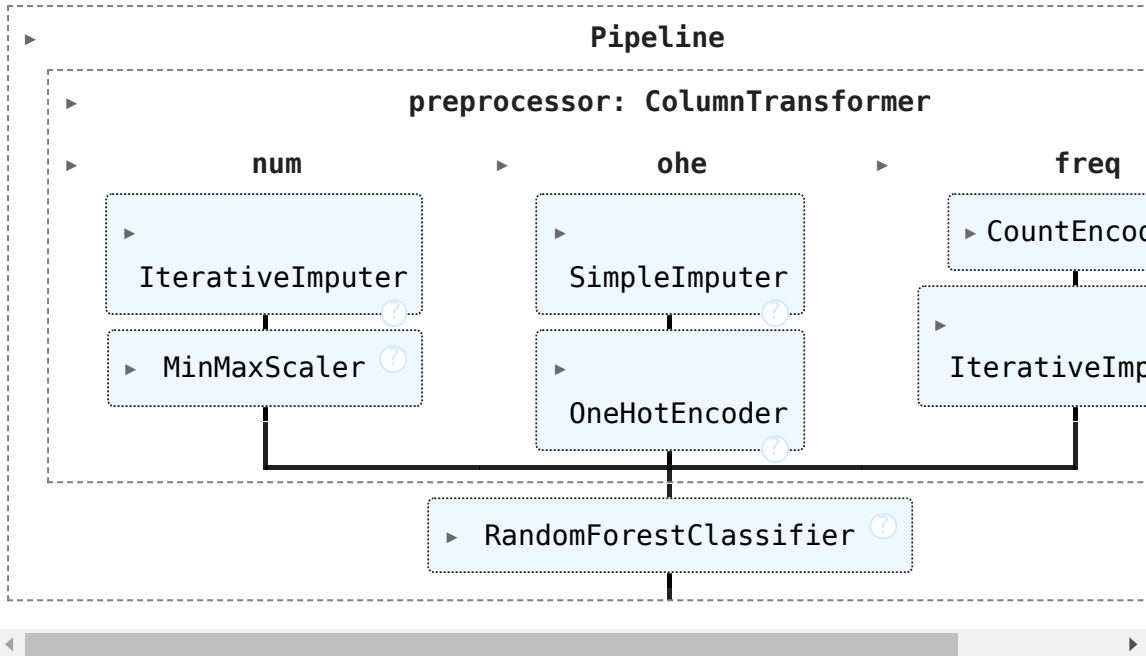
#Print the best parameters for the model
print(output_rfc.best_params_)

# Retrieve the best estimator from GridSearchCV
best_model_rfc=output_rfc.best_estimator_
```

```
{'classifier__class_weight': 'balanced_subsample', 'classifier__max_depth': 5, 'classifier__min_samples_leaf': 1, 'classifier__n_estimators': 100}
```

```
In [ ]: #Fit the best parameters to the training data
output_rfc.best_estimator_.fit(X_tr, y_tr)
```


Out[]:




```
In [ ]: # Evaluating various metrics of the random forest classifier
evaluate(best_model_rfc, X_tr, X_te, y_tr, y_te)
```

Training Scores:
Train Accuracy: 0.7930610812075825
Train Precision: 0.5088612368024132
Train Recall: 0.7433213990636188
Train F1-Score: 0.604141018466704

Testing Scores:
Test Accuracy: 0.7954598642639832
Test Precision: 0.5126865671641792
Test Recall: 0.7566079295154186
Test F1-Score: 0.6112099644128114



```
In [ ]: RocCurveDisplay.from_estimator(best_model_rfc, X_te, y_te, name='random f
plt.show())
```



This random forest classifier model also has low precision and fl scores. It has an AUC score of 0.85, which is slightly better than the decision tree model above. This model is not overfitting to a great extent.

MODEL 4: kNN Classifier

```
In [ ]: kNN = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', KNeighborsClassifier())
])
```

```
In [ ]: cross_validate(kNN, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([4.47753859, 2.20786595, 2.81130219, 2.70128775, 2.59
447789]),
'score_time': array([0.18522739, 0.15040016, 0.15812731, 0.16451311, 0.
19717407]),
'test_score': array([0.78911963, 0.79584674, 0.79256875, 0.79286132, 0.
79724985]),
'train_score': array([0.83924523, 0.84056169, 0.84320608, 0.84050022,
0.84064648])}
```

```
In [ ]: # Define your grid for RandomizedSearchCV
param_grid = {
    "classifier_n_neighbors": [5, 10],
    "classifier_weights": ['uniform', 'distance'],
    "classifier_p": [1, 2, 3],
}
```

```
# Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(knn,
                    param_grid,
                    scoring=['f1','precision'],
                    refit = 'f1')

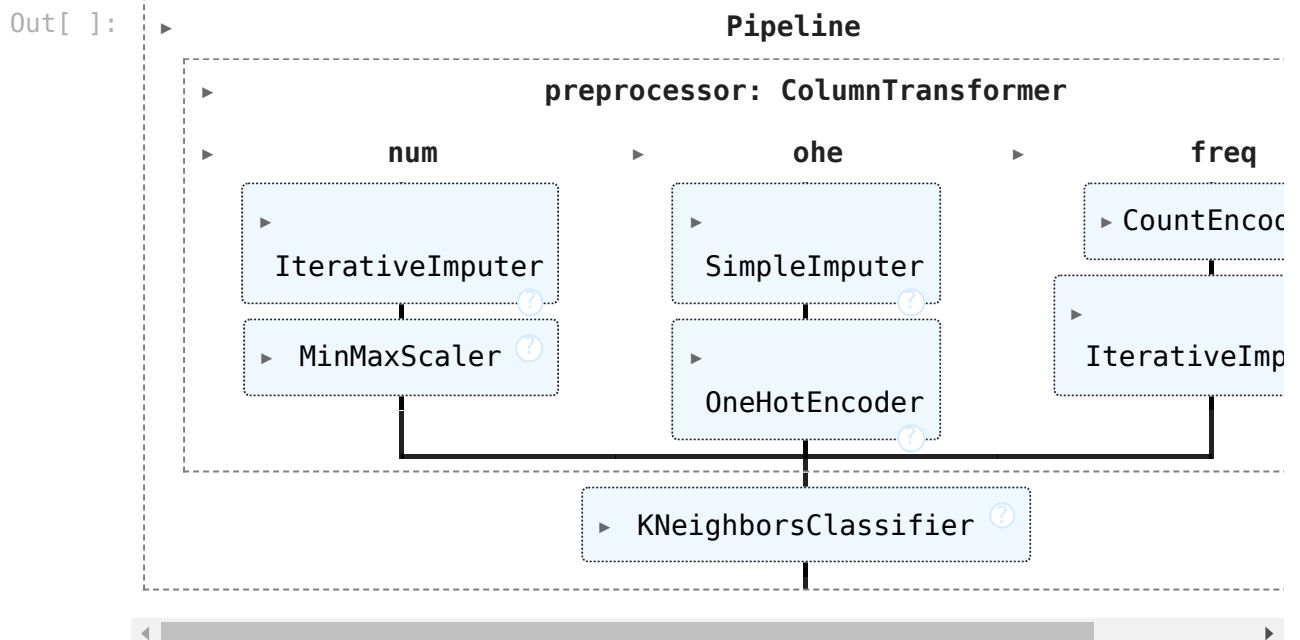
#fit the GridSearchCV
output_knn = grid.fit(X_tr, y_tr)

#Print the best parameters for the model
print(output_knn.best_params_)

# Retrieve the best estimator from GridSearchCV
best_model_KNN =output_knn.best_estimator_
```

```
{'classifier__n_neighbors': 5, 'classifier__p': 1, 'classifier__weights':
'distance'}
```

```
In [ ]: output_knn.best_estimator_.fit(X_tr, y_tr)
```



```
In [ ]: # evaluating the KNN model for various metrics
evaluate(best_model_KNN, X_tr, X_te, y_tr, y_te)
```

Training Scores:

Train Accuracy: 1.0

Train Precision: 1.0

Train Recall: 1.0

Train F1-Score: 1.0


Testing Scores:

Test Accuracy: 0.8013105546454482

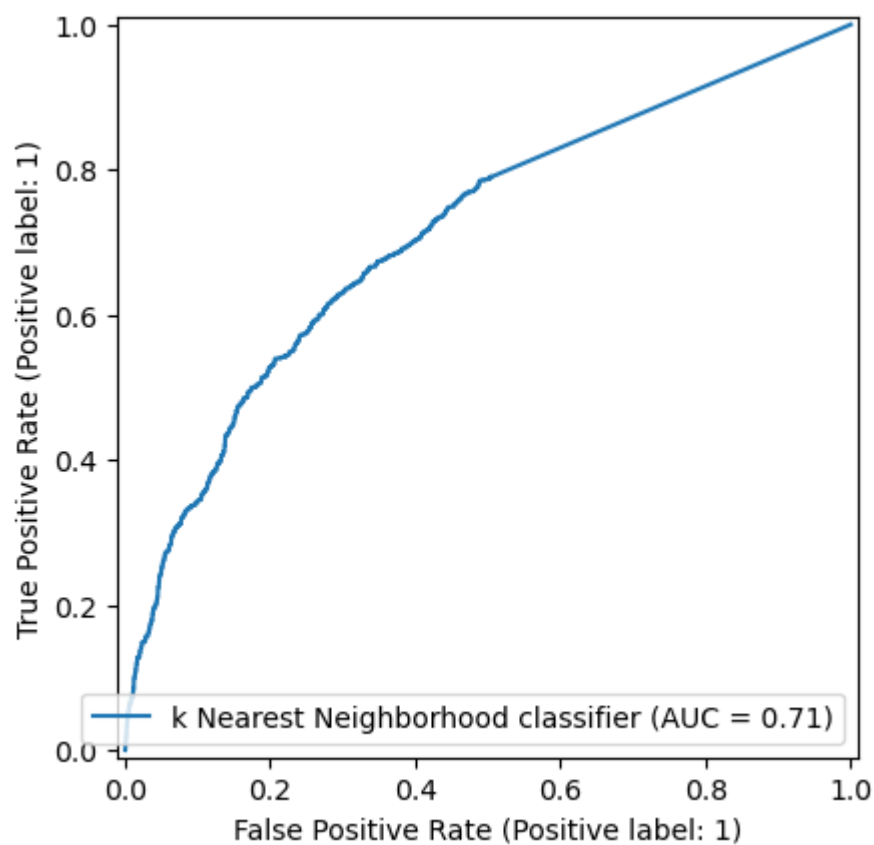
Test Precision: 0.5700712589073634

Test Recall: 0.2643171806167401

Test F1-Score: 0.3611738148984199



```
In [ ]: # Plotting the roc-auc curve for the KNN model
RocCurveDisplay.from_estimator(best_model_KNN, X_te, y_te, name='k Neares
plt.show()
```



This KNN model is definitely overfitting; the training data has perfect scores for all metrics, whereas the testing data scores are much lower. The AUC score is also much

lower than on previous models.

MODEL 5: Gradient Boosting Classifier

```
In [ ]: gbc = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', GradientBoostingClassifier(random_state=42))
])
```

```
In [ ]: cross_validate(gbc, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([6.09436631, 5.15299273, 5.54063416, 5.6029582 , 5.42
427301]),
'score_time': array([0.0781045 , 0.07004285, 0.08713555, 0.07730174, 0.
07750106]),
'test_score': array([0.85288096, 0.85434338, 0.85225278, 0.85137507, 0.
85868929]),
'train_score': array([0.86338038, 0.86506253, 0.86192775, 0.86302472,
0.86295159])}
```

```
In [ ]: # Let's do GridSearchCV
param_grid = {
    "classifier__n_estimators": [100, 200],
    "classifier__max_depth" : [1, 2, 5],
    "classifier__learning_rate": [1, 0.1, 0.01],
}

# Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(gbc,
                    param_grid,
                    scoring=['f1', 'precision'],
                    refit = 'f1') # 2*3*3*5 for CV

# Fit GridSearchCV
output_gbc = grid.fit(X_tr, y_tr)

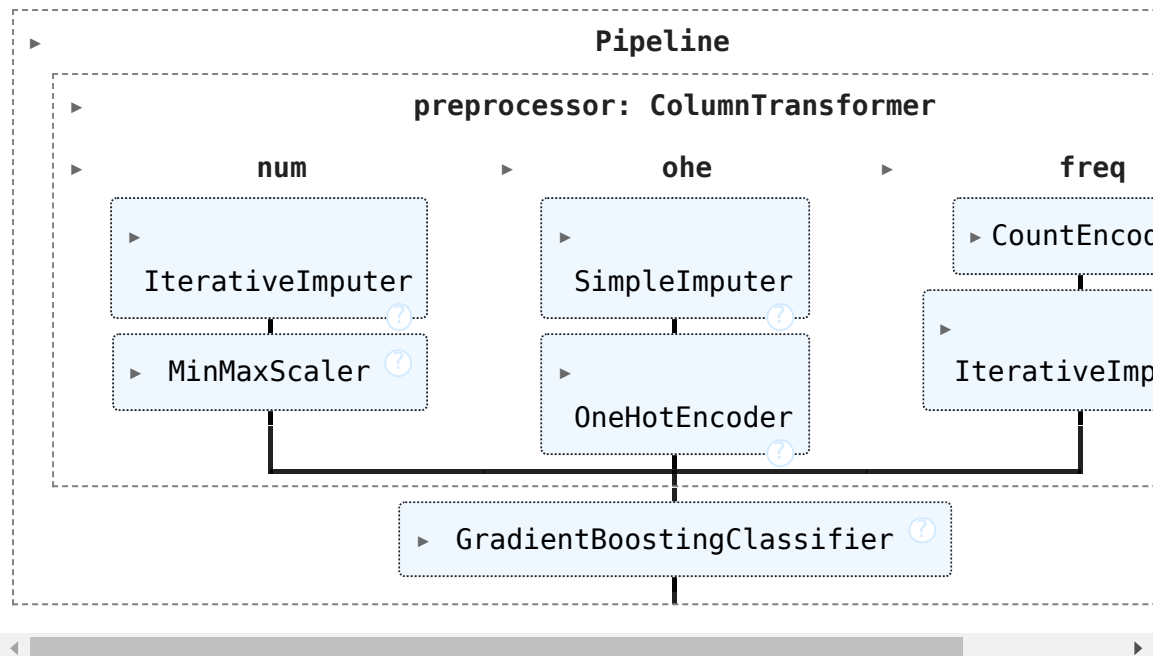
# Print the best parameters found by GridSearchCV
print("Best parameters found: ", output_gbc.best_params_)

# Retrieve the best estimator from GridSearchCV
best_model_gbc = output_gbc.best_estimator_
```

Best parameters found: {'classifier__learning_rate': 1, 'classifier__max_depth': 2, 'classifier__n_estimators': 100}

```
In [ ]: # Fit the best estimator on the training data
best_model_gbc.fit(X_tr, y_tr)
```

Out[]:



```
In [ ]: # Evaluate the best model using the evaluate function
evaluate(best_model_gbc, X_tr, X_te, y_tr, y_te)
```

Training Scores:

Train Accuracy: 0.8706412356658085

Train Precision: 0.772239263803681

Train Recall: 0.5546681354998623

Train F1-Score: 0.6456162846610034


Testing Scores:

Test Accuracy: 0.847648022466651


Test Precision: 0.6843615494978479

Test Recall: 0.525330396475771

Test F1-Score: 0.594392523364486



```
In [ ]: # Plotting the roc-auc curve for the KNN model
RocCurveDisplay.from_estimator(best_model_gbc, X_te, y_te, name='k Neares
plt.show()
```



MODEL 6: XG Boosting Classifier

```
In [ ]: xgb = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', xgboost.XGBClassifier(random_state=42))
])
```

```
In [ ]: cross_validate(xgb, X_tr, y_tr, return_train_score=True)
```

```
Out[ ]: {'fit_time': array([5.63627338, 2.99075842, 3.72523427, 3.32572556, 3.25
237846]),
'score_time': array([0.09734845, 0.06626916, 0.08217311, 0.08071756, 0.
07704997]),
'test_score': array([0.84235156, 0.84586136, 0.84435342, 0.84669397, 0.
85342305]),
'train_score': array([0.95370438, 0.95260733, 0.95443908, 0.9510019 ,
0.95575545])}
```

```
In [ ]: # Let's do GridSearchCV
param_grid = {
    "classifier__n_estimators": [100, 200],
    "classifier__max_depth" : [1, 2],
    "classifier__learning_rate": [1, 0.1],
}

# Setup GridSearchCV with multiple scoring metrics
grid = GridSearchCV(xgb,
```



```

        param_grid,
        scoring=['f1','precision'],
        refit = 'f1')

# Fit GridSearchCV
output_xgb = grid.fit(X_tr, y_tr)

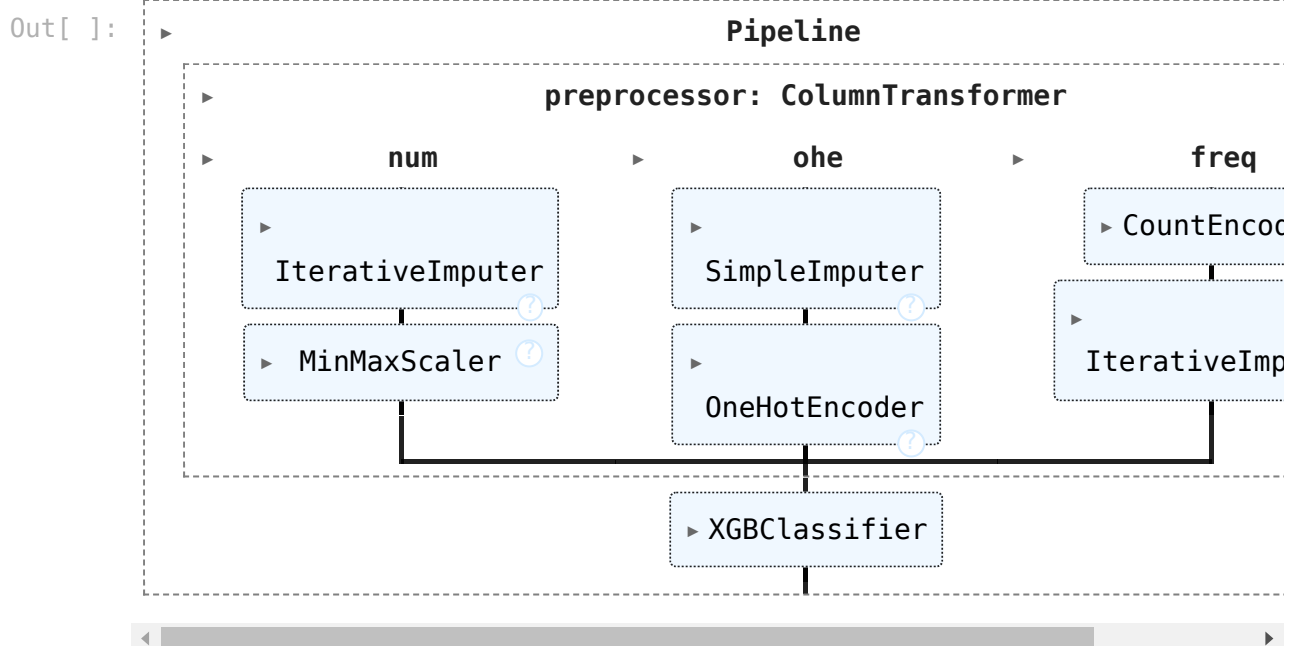
#Print the best parameters of the model
print(output_xgb.best_params_)

# Retrieve the best estimator from GridSearchCV
best_model_xgb=output_xgb.best_estimator_

```

```
{'classifier__learning_rate': 1, 'classifier__max_depth': 2, 'classifier__
n_estimators': 100}
```

```
In [ ]: c#Fit the best model to our training data
best_model_xgb.fit(X_tr ,y_tr)
```



```
In [ ]: # Evaluate the best model using the evaluate function
evaluate(best_model_xgb, X_tr, X_te, y_tr, y_te)
```

Training Scores:

Train Accuracy: 0.8670138076293002

Train Precision: 0.754688672168042

Train Recall: 0.554117323051501

Train F1-Score: 0.6390344608543751


Testing Scores:

Test Accuracy: 0.8481160776971682


Test Precision: 0.6842105263157895

Test Recall: 0.5297356828193832

Test F1-Score: 0.5971446306641838



```
In [ ]: # Plotting the roc-auc curve for the XGB model
RocCurveDisplay.from_estimator(best_model_xgb, X_te, y_te, name='XGBoost')
plt.show()
```



This model gave us similar scores to the gradient boosting model, but the gradient boosting model has the best AUC score and precision score. So, we will choose the

gradient boosting classifier as the final model.

Comparison of Model ROC Curves


```
In [ ]: # Function to plot ROC curves for multiple models
def plot_roc_curves(models, X_test, y_test):
    fig, ax = plt.subplots(figsize=(10, 10))

    for model in models:
        RocCurveDisplay.from_estimator(model['estimator'], X_test, y_test)

    plt.show()

# Define your models
models = [
    {'estimator': dummy_model, 'name': 'Dummy model'},
    {'estimator': best_model_KNN, 'name': 'k Nearest Neighborhood classif'},
    {'estimator': best_model_logreg, 'name': 'Logistic regression'},
    {'estimator': best_model_dtc, 'name': 'Decision tree classification'},
    {'estimator': best_model_rfc, 'name': 'Random forest classifier'},
    {'estimator': best_model_gbc, 'name': 'Gradient boosting classifier'},
    {'estimator': best_model_xgb, 'name': 'XGBoost Classifier'}
]

# Plot ROC curves
plot_roc_curves(models, X_te, y_te)
```



'Final' Model: Gradient Boosting Classifier

From the ROC curves comparison above, and confusion matrix of the method, we decided to choose Gradient Boosting Classifier as our final model.

```
In [ ]: final_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', GradientBoostingClassifier(learning_rate=0.1, n_estimators=100,
                                             max_depth=5,
                                             random_state=42))
])
output_final_model = final_model.fit(X_tr, y_tr)
evaluate(final_model, X_tr, X_hold, y_tr, y_hold)
```

Training Scores:

Train Accuracy: 0.9002457289960215

Train Precision: 0.8543046357615894

Train Recall: 0.6394932525475076

Train F1-Score: 0.731453772247598


Testing Scores:

Test Accuracy: 0.845376263571696


Test Precision: 0.6841477949940405

Test Recall: 0.505726872246696

Test F1-Score: 0.5815602836879432




```
In [ ]: # Plotting the roc-auc curve for the XGB model
RocCurveDisplay.from_estimator(final_model, X_te, y_te, name='Gradient Bo
plt.show()
```



```
In [ ]: # we can calculate predictions for confusion matrix
hold_preds_final_model = final_model.predict(X_hold)
cm=confusion_matrix(y_hold, hold_preds_final_model, normalize='true')
```

```
In [ ]: # Plot heatmap for final model's confusion matrix for better visualization
ax= plt.subplot()
sns.heatmap(cm, annot=True, ax=ax, cbar=False, linewidths=1, cmap="YlGnBu")

ax.set_title('Confusion Matrix')
ax.set_xlabel('Predictions', fontsize = 15)
ax.set_ylabel('Actuals', fontsize = 15)
ax.xaxis.set_ticklabels(['Not Vaccinated', 'Vaccinated'])
ax.yaxis.set_ticklabels(['Not Vaccinated', 'Vaccinated'], rotation = 0, f
```



Evaluation

Our baseline model had an accuracy score of 78%, but a score of zero for precision, recall, and f1 scores. When we compare all of our following models to this baseline, all have much better precision, recall, and f1 scores, and many have higher accuracy scores. The decision tree model is not overfitting, but it has a low precision score, as well as a low f1 score. However, it has an AUC score of 0.84, which is fairly high, meaning that it does an adequate job of maximizing true positives and minimizing the false positives. The decision tree model is not overfitting. This logistic regression model has low precision and f1 scores, and has an AUC equal to the decision tree AUC above. This model is not overfitting. The random forest classifier model also has low precision and f1 scores. It has an AUC score of 0.85, which is slightly better than the decision tree model above. This model is not overfitting to a great extent. The KNN model is definitely overfitting; the training data has perfect scores for all metrics, whereas the testing data scores are much lower. The AUC score is also much lower than on previous models. The Gradient Boosting model has the highest overall scores of all the models done so far, and also does the best job of minimizing the false positives. This is our candidate for the final model. The XG Boost model gave us similar scores to the gradient boosting model, but the gradient boosting model has the best AUC score and precision score. So, we will choose the gradient boosting classifier as the final model.

Our final model did not overfit to the training set, we got similar AUC, precision and accuracy scores for the holdout set. Because the model does a good job of minimizing the false positive rate on the hold out data, we are fairly confident that it will generalize well to unseen data and will accurately help public health officials determine

the people who didn't get the vaccine. We are going to look into feature importances to understand the relationship between the features and vaccination behavior.